What You Must Know about Memory, Caches, and Shared Memory

Kenjiro Taura

Contents

Contents

Introduction

- so far, we have learned
  - parallelization across cores,
  - vectorization (SIMD) within a core, and
  - instruction level parallelism
- another critical factor you must know to understand program performance is *data access*

Why data access is so important?

- no data, no computation

```
1   for (k = 0; k < A.nnz; k++) {
2     i,j,Aij = A.elems[k];
3     y[i] += Aij * x[j];
4   }
```

```
1   for (i = 0; i < M; i++)
2     for (j = 0; j < N; j++)
3       for (k = 0; k < K; k++)
4         C(i,j) += A(i,k) * B(k,j);
```

Why data access is so important?

- no data, no computation

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

```
1  for (i = 0; i < M; i++)
2    for (j = 0; j < N; j++)
3      for (k = 0; k < K; k++)
4        C(i,j) += A(i,k) * B(k,j);
```

- accessing data is sometimes *far more costly* than calculation

Why data access is so important?

- no data, no computation

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

```
1  for (i = 0; i < M; i++)
2    for (j = 0; j < N; j++)
3      for (k = 0; k < K; k++)
4        C(i,j) += A(i,k) * B(k,j);
```

- accessing data is sometimes *far more costly* than calculation
- moreover, the cost of the same data access instruction
  significantly differs depending on *where dare are coming from*
  - registers
  - caches
  - main memory
  - another processor's cache

Conceptual goals of the study

- understand how are processors, caches and memory are connected
- understand the behavior of caches, so as to reason about how much traffic the algorithm will generate between main memory ↔ caches (and among cache levels)
- ⇒ be able to reason about *a performance limit of your program, due to the memory*

Pragmatic goals of the study

- latency: get a sense of how many cycles it takes to get data from main memory and caches

Pragmatic goals of the study

- latency: get a sense of how many cycles it takes to get data from main memory and caches
- bandwidth: get a sense of how much data CPU can bring from main memory and caches

Pragmatic goals of the study

- latency: get a sense of how many cycles it takes to get data from main memory and caches
- bandwidth: get a sense of how much data CPU can bring from main memory and caches
- what does "memory bandwidth" we see in a processor spec sheet really mean? e.g.,
  - the processor data sheet of E5-2698 (68 GB/s):
    `http://ark.intel.com/products/81060/Intel-Xeon-Processor-E5-2698-v3-40M-Cache-2_30-GHz`
  - in general,

    8 bytes × DDR frequency × memory channel, per CPU socket
  - our CPU (Ice Lake Xeon Platinum 8368)

    8 bytes × 3200 MHz × 8 channels ≈ 200 GB/sec per socket

    200 × 2 sockets = 400 GB/sec in the entire node

Pragmatic goals of the study

- latency: get a sense of how many cycles it takes to get data from main memory and caches
- bandwidth: get a sense of how much data CPU can bring from main memory and caches
- what does "memory bandwidth" we see in a processor spec sheet really mean? e.g.,
  - the processor data sheet of E5-2698 (68 GB/s):
    http://ark.intel.com/products/81060/Intel-Xeon-Processor-E5-2698-v3-40M-Cache-2_30-GHz
  - in general,

    8 bytes × DDR frequency × memory channel, per CPU socket

  - our CPU (Ice Lake Xeon Platinum 8368)

    8 bytes × 3200 MHz × 8 channels ≈ 200 GB/sec per socket

    200 × 2 sockets = 400 GB/sec in the entire node

- Can we achieve this easily? If not, when/how can we?

Contents

What does memory performance imply for FLOPS?

- many computationally *efficient* algorithms do not touch the same data too many times

What does memory performance imply for FLOPS?

- many computationally *efficient* algorithms do not touch the same data too many times
- e.g., $O(n)$ algorithms $\rightarrow$ uses a single element only a constant number of times (on average)

What does memory performance imply for FLOPS?

- many computationally *efficient* algorithms do not touch the same data too many times
- e.g., $O(n)$ algorithms $\rightarrow$ uses a single element only a constant number of times (on average)
- if data $\gg$ cache for such an algorithm, the algorithm's performance is often limited by the memory bandwidth (or, worse, latency), *not processor's compute throughput*

Example: SpMV

- remember COO

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

Example: SpMV

- remember COO

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

- accesses 16 nnz bytes and performs 2 nnz flops
  - assuming elements of `double` (8 bytes) and indexes of `int`s
    (4 bytes × 2), not counting access to `x` and `y`
  - details aside, it performs only *an FMA / element*

Example: SpMV

- remember COO

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

- accesses 16 nnz bytes and performs 2 nnz flops
  - assuming elements of `double` (8 bytes) and indexes of `int`s (4 bytes × 2), not counting access to `x` and `y`
  - details aside, it performs only *an FMA / element*
- to achieve Skylake-X peak (16 DP FMAs per core per cycle), a core must access 16 matrix elements (= 256 bytes) / cycle

Example: SpMV

- remember COO

```
1  for (k = 0; k < A.nnz; k++) {
2    i,j,Aij = A.elems[k];
3    y[i] += Aij * x[j];
4  }
```

- accesses 16 nnz bytes and performs 2 nnz flops
  - assuming elements of double (8 bytes) and indexes of ints
    (4 bytes × 2), not counting access to x and y
  - details aside, it performs only *an FMA / element*
- to achieve Skylake-X peak (16 DP FMAs per core per cycle),
  a core must access 16 matrix elements (= 256 bytes) / cycle

Memory-bound algorithms (applications)

- say an algorithm performs C flops (or computation in more general) on N bytes of data
  - assume it needs to access every element of the N bytes at least once (likely the case)

Memory-bound algorithms (applications)

- say an algorithm performs C flops (or computation in more general) on N bytes of data
  - assume it needs to access every element of the N bytes at least once (likely the case)
- there are two obvious lower bounds on the time to complete the algorithm

$$T \geq \frac{C}{\text{the peak FLOPS}} \quad \text{(compute)}$$

$$T \geq \frac{N}{\text{the peak memory bandwidth}} \quad \text{(memory)}$$

Memory-bound algorithms (applications)

- say an algorithm performs C flops (or computation in more general) on N bytes of data
  - assume it needs to access every element of the N bytes at least once (likely the case)
- there are two obvious lower bounds on the time to complete the algorithm

$$T \geq \frac{C}{\text{the peak FLOPS}} \quad \text{(compute)}$$

$$T \geq \frac{N}{\text{the peak memory bandwidth}} \quad \text{(memory)}$$

- often, the latter is much larger and such algorithms are called *"memory-bound"*
- $O(N)$, $O(N \log N)$ algorithms are almost always memory bound

Memory-bound algorithms (applications)

- memory-bound $\iff$

$$\frac{C}{\text{the peak FLOPS}} \ll \frac{N}{\text{the peak memory bandwidth}}$$

$\iff$

$$\frac{C}{N} \ll \frac{\text{the peak FLOPS}}{\text{the peak memory bandwidth}}$$

- - the LHS: *arithmetic intensity* or *compute intensity* of the algorithm
  - the reciprocal of RHS: the *byte per FLOPS* of the machine
- note that being memory-bound suggests it is inefficient in the processor utilization view point, but it is efficient in time-complexity sense *(it is not necessarily a bad thing)*

Note: dense matrix-vector multiply

- the same argument applies even if the matrix is *dense*

```
1  for (i = 0; i < M; i++)
2    for (j = 0; j < N; j++)
3      y[i] += a[i][j] * x[j];
```

Note: dense matrix-vector multiply

- the same argument applies even if the matrix is *dense*

```
1   for (i = 0; i < M; i++)
2     for (j = 0; j < N; j++)
3       y[i] += a[i][j] * x[j];
```

- MN flops on $(MN + M + N)$ elements

Note: dense matrix-vector multiply

- the same argument applies even if the matrix is *dense*

```
1   for (i = 0; i < M; i++)
2     for (j = 0; j < N; j++)
3       y[i] += a[i][j] * x[j];
```

- MN flops on $(MN + M + N)$ elements
- $\Rightarrow$ it performs only an FMA / matrix element

Dense matrix-matrix multiply

- the argument does *not* apply to matrix-matrix multiply
  (we've been trying to get close to CPU peak)

$$M \begin{array}{|c|} \hline N \\ C \\ \hline \end{array} \mathrel{+}= \begin{array}{|c|} \hline K \\ A \\ \hline \end{array} * \begin{array}{|c|} \hline N \\ B \\ \hline \end{array} K$$
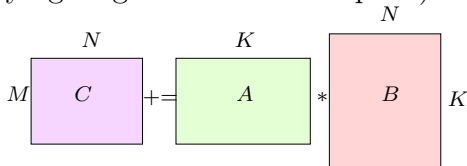
Dense matrix-matrix multiply

- the argument does *not* apply to matrix-matrix multiply
  (we've been trying to get close to CPU peak)
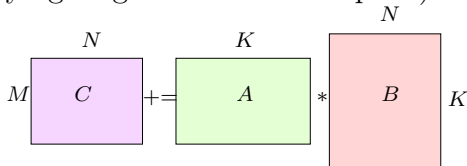
- for $N \times N$ square matrices, it performs N^3 FMAs on $3N^2$
  elements

Why dense matrix-matrix multiply *can* be efficient?

- assume $M \sim N \sim K$

```
1  for (i = 0; i < M; i++)
2    for (j = 0; j < N; j++)
3      for (k = 0; k < K; k++)
4        C(i,j) += A(i,k) * B(k,j);
```

- a microscopic argument
  - the innermost statement

```
1  C(i,j) += A(i,k) * B(k,j)
```

  still performs (only) 1 FMA for accessing 3 elements
  - but the same element (say `C(i,j)`) is used many (K) times in the innermost loop
  - similarly, the same `A(i,k)` is used N times
  - $\Rightarrow$ after you use an element, *if you reuse it many times before it is evicted from a cache (even a register)*, then the memory traffic is hopefully not a bottleneck

A simple `memcpy` experiment ...

```
1  double t0 = cur_time();
2  memcpy(a, b, nb);
3  double t1 = cur_time();
```

A simple `memcpy` experiment . . .

```
1  double t0 = cur_time();
2  memcpy(a, b, nb);
3  double t1 = cur_time();
```

```
1  $ gcc -O3 memcpy.c
2  $ ./a.out $((1 << 26))  # 64M long elements = 512MB
3  536870912 bytes copied in 0.117333 sec 4.575611 GB/sec
```

A simple `memcpy` experiment ...

```
1  double t0 = cur_time();
2  memcpy(a, b, nb);
3  double t1 = cur_time();
```

```
1  $ gcc -O3 memcpy.c
2  $ ./a.out $((1 << 26))  # 64M long elements = 512MB
3  536870912 bytes copied in 0.117333 sec 4.575611 GB/sec
```

- much lower than the advertised number ...

Contents

Cache and memory in a single-core processor

you almost certainly know this (*caches* and main memory), don't you?

recent processors have *multiple levels* of caches (L1, L2, ...)
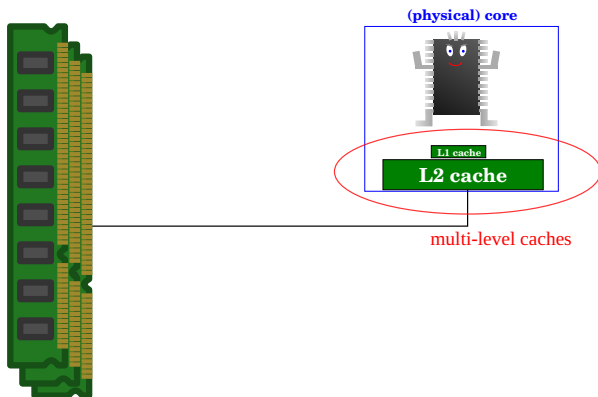
..., with multicores in a chip, ...

- a single chip has several cores
- each core has its *private* caches (typically, L1 and L2)
- cores in a chip share a cache (typical, L3) and main memory

**hardware thread
(virtual core, CPU)**

**memory
controller**

..., with simultaneous multithreading (SMT) in a core, ...

- each core has two *hardware threads*, which share L1/L2 caches and some or all execution units
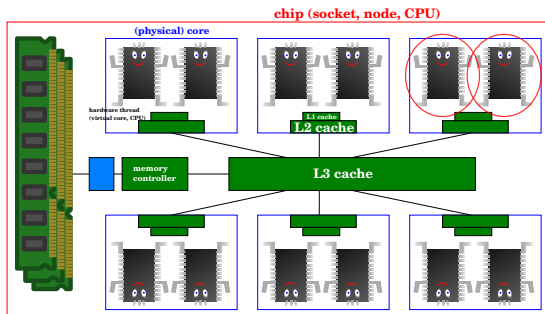
..., and with multiple sockets per node.

- each node has several chips (sockets), connected via an interconnect (e.g., Intel QuickPath, AMD HyperTransport, etc.)
- each socket serves a part of the entire main memory
- each core can still access any part of the entire main memory

Today's typical single compute node

Typical cache sizes

- L1 : 16KB - 64KB/core
- L2 : 256KB - 1MB/core
- L3 : $\sim$ 50MB/socket

Cache 101

- speed :
$$L1 > L2 > L3 > \text{main memory}$$

Cache 101

- speed :
  $$L1 > L2 > L3 > \text{main memory}$$

- capacity :
  $$L1 < L2 < L3 < \text{main memory}$$

Cache 101

- speed :

$$L1 > L2 > L3 > \text{main memory}$$

- capacity :

$$L1 < L2 < L3 < \text{main memory}$$

- each cache holds a subset of data in the main memory

$$L1, L2, L3 \subset \text{main memory}$$

Cache 101

- speed :
$$L1 > L2 > L3 > \text{main memory}$$

- capacity :
$$L1 < L2 < L3 < \text{main memory}$$

- each cache holds a subset of data in the main memory
$$L1, L2, L3 \subset \text{main memory}$$

- typically but not necessarily,
$$L1 \subset L2 \subset L3 \subset \text{main memory}$$

Cache 101

- speed :
$$L1 > L2 > L3 > \text{main memory}$$

- capacity :
$$L1 < L2 < L3 < \text{main memory}$$

- each cache holds a subset of data in the main memory
$$L1, L2, L3 \subset \text{main memory}$$

- typically but not necessarily,
$$L1 \subset L2 \subset L3 \subset \text{main memory}$$

- *which subset is in caches?* $\rightarrow$ cache management (replacement) policy

- a cache generally holds data in *recently accessed* addresses, up to its capacity

Cache management (replacement) policy

- a cache generally holds data in *recently accessed* addresses, up to its capacity
- this is accomplished by the LRU replacement policy (or its approximation):
  - every time a load/store instruction misses a cache, *the least recently used data in the cache will be replaced*

Cache management (replacement) policy

- a cache generally holds data in *recently accessed* addresses, up to its capacity
- this is accomplished by the LRU replacement policy (or its approximation):
  - every time a load/store instruction misses a cache, *the least recently used data in the cache will be replaced*
- $\Rightarrow$ a (very crude) approximation; data in 32KB L1 cache

  $\approx$ *most recently accessed 32K distinct addresses*

Cache management (replacement) policy

- a cache generally holds data in *recently accessed* addresses, up to its capacity
- this is accomplished by the LRU replacement policy (or its approximation):
  - every time a load/store instruction misses a cache, *the least recently used data in the cache will be replaced*
- ⇒ a (very crude) approximation; data in 32KB L1 cache

  $$\approx \textit{most recently accessed 32K distinct addresses}$$

- due to implementation constraints, real caches are slightly more complex

Cache organization : cache line

- a cache = a set of fixed size *lines*
  - typical line size = 64 bytes or 128 bytes,

a 32KB cache with 64 bytes
lines (holds most recently
accessed 512 distinct blocks)

Cache organization : cache line

64 bytes

cache line

512 lines

- a cache = a set of fixed size *lines*
    - typical line size = 64 bytes or 128 bytes,
- a single line is the minimum unit of data transfer between levels (and replacement)

a 32KB cache with 64 bytes lines (holds most recently accessed 512 distinct blocks)

Cache organization : cache line

64 bytes
cache line

512 lines

- a cache = a set of fixed size *lines*
  - typical line size = 64 bytes or 128 bytes,
- a single line is the minimum unit of data transfer between levels (and replacement)

a 32KB cache with 64 bytes lines (holds most recently accessed 512 distinct blocks)

data in 32KB L1 cache (line size 64B)

$\approx$ *most recently accessed 512 distinct lines*

Associativity of caches

full associative: a block can occupy any line in the cache, regardless of its address

direct map: a block has only *one* designated "seat" (*set*), determined by its address

K-way set associative: a block has K designated "seats", determined by its address

- direct map ≡ 1-way set associative

- full associative ≡ ∞-way set associative

An example cache organization

- Ice Lake Platinum 8368

level	line size	capacity	associativity
L1	64B	48KB/core	12
L2	64B	512KB/core?	8
L3	64B	57MB/socket (38 cores)	??

- Skylake-X Gold 6130

level	line size	capacity	associativity
L1	64B	32KB/core	8
L2	64B	1MB/core	16
L3	64B	22MB/socket (16 cores)	11

- Ivy Bridge E5-2650L

level	line size	capacity	associativity
L1	64B	32KB/core	8
L2	64B	256KB/core	8
L3	64B	36MB/socket (8 cores)	20

What you need to remember in practice about associativity

- *avoid having addresses used together "a-large-power-of-two" bytes apart*
- corollaries:
    - avoid having a matrix with a-large-power-of-two number of columns (a common mistake)
    - avoid managing your memory by chunks of large-powers-of-two bytes (a common mistake)
    - avoid experiments only with $n = 2^p$ (a *very* common mistake)
- why? $\Rightarrow$ they tend to go to the same set and "conflict misses" result

Conflict misses

- consider 8-way set associative L1 cache with 32KB (line size = 64B)
    - 32KB/64B = 512 (= 2^9) lines
    - 512/8 = 64 (= 2^6) sets
- $\Rightarrow$ given an address a, $a[6:11]$ (6 bits) designates the set it belongs to (indexing)

address within a line ($2^6 = 64$ bytes)

index the set in the cache (among $2^6 = 64$ sets)

- if two addresses a and b are a multiple of 2^{12} (4096) bytes apart, they go to the same set

A convenient way to understand conflicts

- it's convenient to think of a cache as two dimensional array of lines. e.g. 32KB, 8-way set associative = 64 (sets) × 8 (ways) array of lines

A convenient way to understand conflicts

- formula 1:

$$\text{worst stride} = \frac{\text{cache size}}{\text{associativity}} \text{ bytes}$$

if addresses are this much apart, they go to the same set
  - e.g., 32KB 8-way set associative ⇒ the worst stride = 4096

A convenient way to understand conflicts

- lesser powers of two are significant too; continuing with the same setting (32KB, 8way-set assocative)

stride	the number of sets they are mapped to	utilization
2048	2	1/32
1024	4	1/16
512	8	1/8
256	16	1/4
128	32	1/2
64	64	1

- formula 2: you stride by

    $P \times$ line size   (P divides S)

    $\Rightarrow$ you utilize only $1/P$ of the capacity

- N.B. formula 1 is a special case, with $P = S$

A remark about virtually-indexed vs. physically-indexed caches

- caches typically use *physical* addresses to select the set an address maps to
- so "addresses" I have been talking about are physical addresses, not virtual addresses you can see as pointer values

a | | address within a line ($2^6 = 64$ bytes)
index the set in the cache

- since virtual $\rightarrow$ physical mapping is determined by the OS (based on the availability of physical memory),

    "two virtual addresses 2^b bytes apart"

    does *not* necessarily imply

    "their physical addresses 2^b bytes apart"

- so what's the significance of the stories so far?

A remark about virtually-indexed vs. physically-indexed caches

- virtual $\rightarrow$ physical translation happens with page granularity (typically, $2^{12} = 4096$ bytes)
- $\rightarrow$ the last 12 bits are intact with the translation

A remark about virtually-indexed vs. physically-indexed caches

- therefore,

  "two virtual *addresses* 2^b *bytes apart"* $\rightarrow$ *"their* physical *addresses* 2^b *bytes apart"*

  for up to page size ($2^b \leq$ page size)

- $\rightarrow$ the formula 2 is valid for strides up to page size

stride	utilization
4096	1/64
2048	1/32
1024	1/16
512	1/8
256	1/4
128	1/2
64	1

changed by address translation

intact with address translation

15 14 12 11 6 5 0

a

256KB/8way

address within a line ($2^6 = 64$

index the set in the cache (among $2^9 = 5$

Remarks applied to different cache levels

stride	utilization
...	$\sim 1/64$
16384	$\sim 1/64$
8192	$\sim 1/64$
4096	$1/64$
2048	$1/32$
1024	$1/16$
512	$1/8$
256	$1/4$
128	$1/2$
64	1

- small caches that use only the last 12 bits to index the set make no difference between virtually- and physically-indexed caches

- for larger caches, the utilization will similarly drop up to stride = 4096, after which it will stay around 1/64

- L1 (32KB/8-way) vs. L2 (256KB/8-way)

Avoiding conflict misses

- e.g., if you have a matrix:

```
1  float a[100][1024];
```

  then `a[i][j]` and `a[i+1][j]` go to the same set in L1 cache;
- ⇒ scanning a column of such a matrix will experience almost 100% cache miss
- avoid it by:

```
1  float a[100][1024+16];
```

What are in the cache?

- consider a cache of
  - capacity $= C$ bytes
  - line size $= Z$ bytes
  - associativity $= K$

What are in the cache?

- consider a cache of
  - capacity $= C$ bytes
  - line size $= Z$ bytes
  - associativity $= K$
- approximation 0.0 (only consider C; $\equiv Z = 1, K = \infty$):

  Cache $\approx$ most recently accessed C distinct addresses

What are in the cache?

- consider a cache of
  - capacity $= C$ bytes
  - line size $= Z$ bytes
  - associativity $= K$
- approximation 0.0 (only consider C; $\equiv Z = 1, K = \infty$):

  Cache $\approx$ most recently accessed C distinct addresses

- approximation 1.0 (only consider C and Z; $K = \infty$):

  Cache $\approx$ most recently accessed C/Z distinct lines

What are in the cache?

- consider a cache of
  - capacity $= C$ bytes
  - line size $= Z$ bytes
  - associativity $= K$
- approximation 0.0 (only consider C; $\equiv Z = 1, K = \infty$):

  Cache $\approx$ most recently accessed C distinct addresses

- approximation 1.0 (only consider C and Z; $K = \infty$):

  Cache $\approx$ most recently accessed C/Z distinct lines

- approximation 2.0 (consider associativity too):
  - depending on the stride of the addresses you use, reason about the utilization (effective size) of the cache
  - in practice, avoid strides of "line size $\times 2^b$"

Contents

Assessing the cost of data access

- we like to obtain cost to access data in each level of the caches as well as main memory
- latency: time until the result of a load instruction becomes available
- bandwidth: the maximum amount of data per unit time that can be transferred between the layer in question to CPU (registers)

Contents

How to measure a latency?

- prepare an array of N records and access them repeatedly

How to measure a latency?

- prepare an array of N records and access them repeatedly
- to measure the *latency*, make sure N load instructions *make a chain of dependencies* (link list traversal)

```
1  for (N times) {
2    p = p->next;
3  }
```

How to measure a latency?

- prepare an array of N records and access them repeatedly
- to measure the *latency*, make sure N load instructions *make a chain of dependencies* (link list traversal)

```
1  for (N times) {
2    p = p->next;
3  }
```

- make sure p->next links all the elements in a random order (the reason becomes clear later)

next pointers
(link all elements in a random order)

cache line size

N elements

Data size vs. latency

- main memory is local to the accessing thread

```
1  $ numactl --cpunodebind 0 --interleave 0 ./mem
2  $ numactl -N 0 -i 0 ./mem   # abbreviation
```


latency per load in a random list traversal [0,1073741824]

How long are latencies

- heavily depends on in which level of the cache data fit
- environment: Skylake-X Xeon Gold 6130 (32KB/1MB/22MB)

size	level	latency (cycles)	latency (ns)
12,736	L1	4.004	1.31
103,616	L2	13.80	4.16
2,964,928	L3	77.40	24.24
301,307,584	main	377.60	115.45

latency per load in a random list traversal [0,1073741824]

A remark about replacement policy

- if a cache stricly follows the LRU replacement policy, once data overflow the cache, repeated access to the data will quickly become *almost-always-miss*
- the "cliffs" in the experimental data look gentler than the theory would suggest

A remark about replacement policy

- if a cache stricly follows the LRU replacement policy, once data overflow the cache, repeated access to the data will quickly become *almost-always-miss*
- the "cliffs" in the experimental data look gentler than the theory would suggest

size to repeatedly scan

A remark about replacement policy

- if a cache stricly follows the LRU replacement policy, once data overflow the cache, repeated access to the data will quickly become *almost-always-miss*
- the "cliffs" in the experimental data look gentler than the theory would suggest

A remark about replacement policy

- part of the gap is due to virtual → physical address translation
- another factor, especially for L3 cache, will be a recent replacement policy for cyclic accesses (c.f. `http://blog.stuffedcow.net/2013/01/ivb-cache-replacement/`)

Latency to a remote main memory

- make main memory remote to the accessing thread

```
1  $ numactl -N 0 -i 1 ./mem
```


latency per load in a random list traversal [0,1073741824]

Contents

Bandwidth of a random link list traversal

$$\text{bandwidth} = \frac{\text{total bytes read}}{\text{elapsed time}}$$

- in this experiment, we set record size = 64

bandwidth of list traversal [0,1073741824]

The "main memory" bandwidth

bandwidth of list traversal [33554432,1073741824]

- $\ll$ the `memcpy` bandwidth we have seen (≈ 4.5 GB/s)
- not to mention the "memory bandwidth" in the spec

Why is the bandwidth so low?

- while traversing a single link list, only a single record access (64 bytes) is "in flight" at a time

(link all elements in a random order)

cache line size

N elements

(physical) core

cache

- in this condition,

How to get more bandwidth?

- just like flops/clock, the only way to get a better throughput (bandwidth) is to perform *many load operations concurrently*

How to get more bandwidth?

- just like flops/clock, the only way to get a better throughput (bandwidth) is to perform *many load operations concurrently*

- there are several ways to make it happen; let's look at conceptually the most straightforward: traverse multiple lists

```
1  for (N times) {
2    p1 = p1->next;
3    p2 = p2->next;
4        ...
5  }
```

Contents

The number of lists vs. bandwidth

bandwidth with a number of chains [0,1073741824]

- let's zoom into "main memory" regime (size > 100MB)

Bandwidth to the local main memory (not cache)

- an almost proportional improvement up to ∼ 10 lists


bandwidth with a number of chains [33554432,1073741824]

Bandwidth to a remote main memory (not cache)

- pattern is the same (improve up to ~ 10 lists)
- remember the remote latency is longer, so the bandwidth is accordingly lower

bandwidth with a number of chains [33554432,1073741824]

The number of lists vs. bandwidth

- observation: bandwidth increase fairly proportionally to the number of lists, matching our understanding, . . .

The number of lists vs. bandwidth

- observation: bandwidth increase fairly proportionally to the number of lists, matching our understanding, . . .

- question: . . . but up to ~ 10, why?

The number of lists vs. bandwidth

- observation: bandwidth increase fairly proportionally to the number of lists, matching our understanding, ...

- question: ... but up to ~ 10, why?
- answer: there is a limit in the number of load operations in flight at a time

Line Fill Buffer

- *Line fill buffer (LFB)* is the processor resource that keeps track of outstanding cache misses, and its size is 10 in Haswell
  - I could not find the definitive number for Skylake-X, but it will probably be the same

Line Fill Buffer

- *Line fill buffer (LFB)* is the processor resource that keeps track of outstanding cache misses, and its size is 10 in Haswell
  - I could not find the definitive number for Skylake-X, but it will probably be the same
- this gives *the maximum attainable bandwidth per core*

$$\frac{\text{cache line size} \times \text{LFB size}}{\text{latency}}$$

Line Fill Buffer

- *Line fill buffer (LFB)* is the processor resource that keeps track of outstanding cache misses, and its size is 10 in Haswell
  - I could not find the definitive number for Skylake-X, but it will probably be the same
- this gives *the maximum attainable bandwidth per core*

$$\frac{\text{cache line size} \times \text{LFB size}}{\text{latency}}$$

- this is what we've seen (still much lower than what we see in the "memory bandwidth" in the spec sheet)

Line Fill Buffer

- *Line fill buffer (LFB)* is the processor resource that keeps track of outstanding cache misses, and its size is 10 in Haswell
  - I could not find the definitive number for Skylake-X, but it will probably be the same
- this gives *the maximum attainable bandwidth per core*

$$\frac{\text{cache line size} \times \text{LFB size}}{\text{latency}}$$

- this is what we've seen (still much lower than what we see in the "memory bandwidth" in the spec sheet)
- how can we go beyond this? $\Rightarrow$ the only way is to *use multiple cores* (covered later)

Contents

Other ways to get more bandwidth

- we've learned:
  - maximum bandwidth ≈ as many memory accesses as possible always in flight
  - there is a limit due to LFB entries (10 in Haswell)

Other ways to get more bandwidth

- we've learned:
    - maximum bandwidth ≈ as many memory accesses as possible always in flight
    - there is a limit due to LFB entries (10 in Haswell)
- so far, we have achieved larger bandwidth by traversing multiple lists explicitly (sometimes difficult if not impossible to apply)

Other ways to get more bandwidth

- we've learned:
    - maximum bandwidth ≈ as many memory accesses as possible always in flight
    - there is a limit due to LFB entries (10 in Haswell)
- so far, we have achieved larger bandwidth by traversing multiple lists explicitly (sometimes difficult if not impossible to apply)
- fortunately, the life is not always that tough; there are other ways to issue many memory accesses concurrently
    1. make addresses sequential
    2. make address generations independent
    3. prefetch by software (make address generations go ahead)
    4. use multiple threads/cores

Other ways to get more bandwidth

- we've learned:
    - maximum bandwidth $\approx$ as many memory accesses as possible always in flight
    - there is a limit due to LFB entries (10 in Haswell)
- so far, we have achieved larger bandwidth by traversing multiple lists explicitly (sometimes difficult if not impossible to apply)
- fortunately, the life is not always that tough; there are other ways to issue many memory accesses concurrently
    1. make addresses sequential
    2. make address generations independent
    3. prefetch by software (make address generations go ahead)
    4. use multiple threads/cores
- remember, all boil down to keep as many memory accesses as possible (up to LFB entries) in flight

Contents

Make addresses sequential

- again build a (single) linked list, but this time, `p->next` always points to the immediately following block
- note that *the instruction sequence is identical* to before; only addresses differ

next pointers
(link all elements in the sequential order)

cache line size

N elements

Bandwidth of traversing address-ordered list

- a factor of 10 faster than random case, but this time with only a single list

bandwidth of random list traversal vs address-ordered list traversal [0,1073741824]

The reason this is faster

- *hardware prefetcher*
- CPU watches the sequence of addresses accessed
- sequential addresses (addresses of a small constant stride) trigger CPU's hardware prefetcher
- CPU issues load instruction ahead of actual data stream on your behalf, to keep the maximum number of loads in flight

next pointers
(link all elements in the sequential order)

cache line size

N elements

Contents

Make address generations independent

- if addresses of memory accesses can be computed without values returned from previous loads, CPU can issue them concurrently

```
1  for (N times) {
2    j = ... /* not use a[·] */
3    a[j];
4  }
```


- note: it's *not* a prefetch (but a real fetch)

Bandwidth when not traversing a list

- ptrchase : chase pointers of a random list
- random : access random addresses, but w/o pointer chasing
- sequential : access sequential addresses, w/o pointer chasing

list traversal vs random access vs sequential access [0,1073741824]

Main memory bandwidth

- pointer chase ≪ random < sequential
- random is ≈ 5x faster than traversing a single random list

list traversal vs random access vs sequential access [33554432,1073741824]

Main memory bandwidth (random vs. sequential)

- sequential gets $\approx$ 3x more bandwidth than random
- may not be as bad as you thought?
- but why is there *any* difference, if both have the same number of loads in flight?

list traversal vs random access vs sequential access [33554432,1073741824]

Random (index) vs. sequential

- if both can have up to 10 (LFB entries) outstanding L1 cache misses, why is there *any* diffference?
- I don't have a definitive answer, but presumably,
  - the hardware prefetcher happens at multiple levels ($\rightarrow$ L1 and $\rightarrow$ L2)
  - prefetchers to L2 are not subject of the LFP entries limit (the limit will be slightly more)
  - prefething to L2 make effective latency to the processor smaller

When "random access" is really bad

- in practice, when random vs. sequential makes a large ($\gg 2$) difference, it's because

$$\text{a single element} < \text{a single cache line}$$

- recall that touching a single byte in a cache line still brings the whole line (64 bytes)
- e.g., if you access an array of `float` (4 bytes) randomly, the bandwidth of *useful* data is amplified by a factor of 16 ($= 64/4$)

Contents

Software prefetch

- hardware prefetch happens only for sequential (a small constant stride) accesses
- for other patterns, you the programmer may know addresses you are going to access soon

Software prefetch

- hardware prefetch happens only for sequential (a small constant stride) accesses
- for other patterns, you the programmer may know addresses you are going to access soon
- *if* you can generate those addresses much ahead of actual load instructions, you can *prefetch* them

Software prefetch

- hardware prefetch happens only for sequential (a small constant stride) accesses
- for other patterns, you the programmer may know addresses you are going to access soon
- *if* you can generate those addresses much ahead of actual load instructions, you can *prefetch* them
- instructions:
  - prefetcht{0,1,2}
  - prefetchnta

Software prefetch

- hardware prefetch happens only for sequential (a small constant stride) accesses
- for other patterns, you the programmer may know addresses you are going to access soon
- *if* you can generate those addresses much ahead of actual load instructions, you can *prefetch* them
- instructions:
  - prefetcht{0,1,2}
  - prefetchnta
- intrinsics:

```
1  __builtin_prefetch(a [, rw, hint ])
```

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful
- why? by the time you can *prefetch* it, you can likewise *load* it!

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful
- why? by the time you can *prefetch* it, you can likewise *load* it!
- in our example,
  - no point in applying it to index-based accesses (CPU will issue many load instructions already)

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful
- why? by the time you can *prefetch* it, you can likewise *load* it!
- in our example,
  - no point in applying it to index-based accesses (CPU will issue many load instructions already)
  - on the other hand, it's difficult to apply it to list traversal (it takes equally long time to generate address to prefetch)

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful
- why? by the time you can *prefetch* it, you can likewise *load* it!
- in our example,
  - no point in applying it to index-based accesses (CPU will issue many load instructions already)
  - on the other hand, it's difficult to apply it to list traversal (it takes equally long time to generate address to prefetch)
- the only way to apply it is to change the data structure of the linked list

How to apply software prefetch?

- truth is, there are actually not many cicumstances this is useful
- why? by the time you can *prefetch* it, you can likewise *load* it!
- in our example,
  - no point in applying it to index-based accesses (CPU will issue many load instructions already)
  - on the other hand, it's difficult to apply it to list traversal (it takes equally long time to generate address to prefetch)
- the only way to apply it is to change the data structure of the linked list
- but how?

How to apply software prefetch?

- have another pointer pointing many elements ahead

```
1  for (N times) {
2    p = p->next;
3    prefetch(p->prefetch);
4  }
```

- it should point to Q elements ahead to have Q concurrent accesses in flight

Result

bandwidth w/ and w/o prefetch [33554432,1073741824]

Summary: bandwidth of various access patterns

- sequential (w/o pointer chase) > sorted list
  > random (w/o pointer chase) ≈ 5 random lists ≈ a random
  list + software prefetch
  > a random list

summary of various access patterns [33554432,1073741824]

Contents

Memory bandwidth with multiple cores

- the bandwidth to a single core is limited by LFB entries and is much lower than the memory bandwidth itself

$$\frac{\text{transfer (line) size } \times \text{ LFB entries}}{\text{latency}}$$

- you can go beyond that by using multiple cores and *this is the only way*

Memory bandwidth with multiple cores

- run up to 16 threads,
- each running on a distinct physical core of a single socket
- allocate all the data on the same socket (`numactl -N 0 -i 0`)
- note: they are still random pointer chasing

bandwidth with a number of threads [33554432,1073741824]
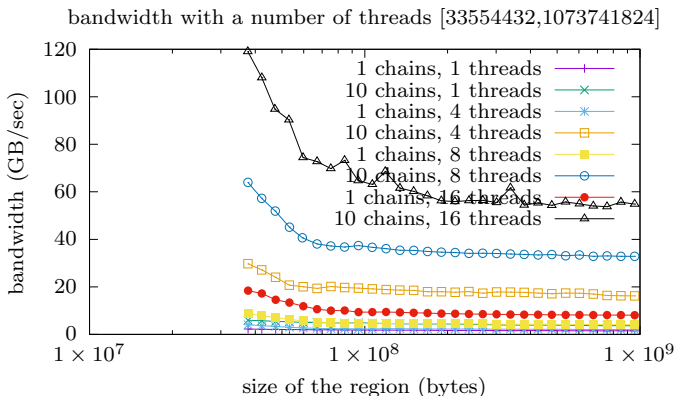
With random indexing and sequential accesses

- similar experiments with random indexing/sequential accesses
- $\sim$ 80 GB/sec with sequential accesses by $\geq$ 12 threads
- the theoretical peak is

  8 bytes $\times$ 2.666 GHz $\times$ 6 channels $= 128$ GB/sec

bandwidth with various methods and number of threads [33554432,1073741824]

With multiple CPU sockets

- the total bandwidth depends on how to place threads and data

threads\data	CPU x	CPU y	all CPUs	local CPU
CPU x	1-local	1-remote	1-all	1-local
all CPUs	all-1	all-1	all-all	all-local

- control threads/data placement by `numactl` command
- combine it with `OMP_PROC_BIND=true` to get a desired effect

numactl command (1)

- usage (see man numactl for details)

```
1  $ numactl options command
```

- for underlying system calls, see man -s 3 numa

- processors
  - -N x runs threads only on the CPU(s) x. e.g.,

  ```
  1  $ numactl -N 0 command    # threads on CPU 0
  ```

  - --physcpubind x runs threads only on *core(s)* x. e.g.,

  ```
  1  # threads on cores 0-11 and 16-27
  2  $ numactl --physcpubind 0-11,16-27 command
  ```

`numactl` command (2)

- memory (data)
  - `-i` y allocates data (physical pages) on CPU(s) y

```
1  $ numactl -i 0,1 command # data on CPU 0 or 1
2  $ numactl -i all command # data on all CPUs
```

  - `-l` allocates physical pages to the CPU that touches the page for the first time (*first touch policy*; the default policy of Linux)

```
1  $ numactl -l command
```

About the `-l` option

- `-l` (equivalent: `--localalloc`) allocates the physical page for a logical page *on the CPU that first touches it (first touch)*
- allocated physical pages do not move thereafter (unless you do so by `move_pages()` system call)
- don't be fooled by its name; it is *not* a policy that automagically makes memory accesses local
- quite contrary, it often makes a *hotspot* in a single CPU, especially when only one thread initializes (first-touches) the data
- `-iall` is not optimal, but often much safer for parallel applications

OpenMP thread placement

- combine them with OMP_NUM_THREADS= and
  OMP_PROC_BIND=true to get a desired effect. e.g.,

```
1  $ OMP_NUM_THREADS=48 OMP_PROC_BIND=true numactl --physcpubind
          0-11,16-27,32-43,48-59 -l command
```

to

- run 12 threads on each CPU (of a host in the big partition)
- and use the first touch policy

Achieved bandwidth

- Skylake X 6130 ×4 CPUs (a host of the "big" partition)
- use 12 (of 16) cores on each CPU
- in each measurement, each thread reads ≈ 640MB sequentially 10 times

setting	threads	bandwidth (GB/sec)
1-local	12	85
1-remote	12	16
1-all	12	57
all-1	48	2
all-all	48	97
all-local	48	320

Remarks on remote access bandwidths

- *numbers for remote accesses* are ridiculously low
- the measurement is repeated 6 times and there were almost no variations in the result (within a few per cents)
- I am suspecting a wrong BIOS snoop setting (`https://software.intel.com/en-us/forums/software-tuning-performance-optimization-platform-mon topic/602160`)

setting	threads	bandwidth (GB/sec)
1-local	12	85
1-remote	12	16
1-all	12	57
all-1	48	2
all-all	48	97
all-local	48	320

Contents

Shared memory

- if thread P writes to an address a and then another thread B reads from a, Q observes the value written by P

Shared memory

- if thread P writes to an address a and then another thread B reads from a, Q observes the value written by P

- ordinary load/store instructions accomplish this *(hardware shared memory)*
- this should not be taken for granted; processors have *caches* and a single address may be cached by multiple cores/sockets

Shared memory

- ⇒ processors sharing memory are running a complex, *cache coherence protocol* to accomplish this
- roughly,

interconnect

Shared memory

- $\Rightarrow$ processors sharing memory are running a complex, *cache coherence protocol* to accomplish this
- roughly,
  1. a write to an address by a processor "invalidates" all other cache lines holding the address, so that no caches hold "stale" values

interconnect

Shared memory

- ⇒ processors sharing memory are running a complex, *cache coherence protocol* to accomplish this
- roughly,
  1. a write to an address by a processor "invalidates" all other cache lines holding the address, so that no caches hold "stale" values
  2. a read to an invalid line causes a miss and searches for a cache holding its "valid" value

interconnect

An example protocol : the MSI protocol

- each line of a cache is in one of the following states

  Modified (■), *Shared* (■), *Invalid* (▬)

An example protocol : the MSI protocol

- each line of a cache is inone of the following states

  Modified (■), *Shared* (■), *Invalid* (▬)

  - Modified (■) ⟺ you can read and write the line without invoking a transaction
  - Shared (■) ⟺ you can read but not write the line without invoking a transaction
  - Invalid (▬) ⟺ you can neither read nor write the line without invoking a transaction

interconnect

An example protocol : the MSI protocol

- a single address may be cached in multiple caches (lines)

interconnect

An example protocol : the MSI protocol

- a single address may be cached in multiple caches (lines)
- ⇒ there are only two legitimate states for each line
  1. one Modified *(owner)* + others Invalid (▬, ▮, ▬, ▬, ▬, . . .)

interconnect

An example protocol : the MSI protocol

- a single address may be cached in multiple caches (lines)
- ⇒ there are only two legitimate states for each line
  1. one Modified *(owner)* + others Invalid (▬, ▮, ▬, ▬, ▬, …)
  2. no Modified (▮, ▬, ▮, ▮, ▬, …)

Cache states and transaction

- suppose a processor reads or writes an address and finds a line caching it
- what happens when the line is in each state:

	Modified	Shared	Invalid
read	hit	hit	read miss
write	hit	write miss	read miss; write miss

Cache states and transaction

- suppose a processor reads or writes an address and finds a line caching it
- what happens when the line is in each state:

	Modified	Shared	Invalid
read	hit	hit	read miss
write	hit	write miss	read miss; write miss

- read miss: $\rightarrow$
  - there may be a cache holding it in Modified state *(owner)*
  - searches for the owner and if found, downgrade it to Shared
  - ▬, ■, ▬, [▬], ▬, ... $\Rightarrow$ ▬, ■, ▬, [■], ▬, ...

Cache states and transaction

- suppose a processor reads or writes an address and finds a line caching it
- what happens when the line is in each state:

	Modified	Shared	Invalid
read	hit	hit	read miss
write	hit	write miss	read miss; write miss

- read miss: →
  - there may be a cache holding it in Modified state *(owner)*
  - searches for the owner and if found, downgrade it to Shared
  - ▬, 🟩, ▬, [▬], ▬, … ⇒ ▬, 🟨, ▬, [🟨], ▬, …
- write miss: →
  - there may be caches holding it in Shared state *(sharer)*
  - searches for sharers and downgrade them to Invalid
  - 🟨, ▬, 🟨, [🟨], ▬, … ⇒ ▬, ▬, ▬, [🟩], ▬, …

MESI and MESIF

- exntensions to MSI have been commonly used

MESI and MESIF

- exntensions to MSI have been commonly used
- MESI: MSI + Exclusive (owned but not modified)
  - when a read request finds no other caches that have the line, it owns it as Exclusive
  - Exclusive lines do not have to be written back to main memory when discarded

MESI and MESIF

- exntensions to MSI have been commonly used
- MESI: MSI + Exclusive (owned but not modified)
  - when a read request finds no other caches that have the line, it owns it as Exclusive
  - Exclusive lines do not have to be written back to main memory when discarded
- MESIF: MESI + Forwarding (a cache responsible for forwarding a line)
  - used in Intel QuickPath
  - when a line is shared by many readers, one is designated as the Forwarder
  - when another cache requests the line, only the forwarder sends it and the new requester becomes the forwarder
  - (in MSI or MESI, all sharers forward it)

How to measure communication latency?

- measure "ping-pong" latency between two threads

```
1  volatile long x = 0;
2  volatile long y = 0;
```

```
1  (ping thread)
2  for (i = 0; i < n; i++) {
3    x = i + 1;
4    while (y <= i) ;
5  }
```

```
1  (pong thread)
2  for (i = 0; i < n; i++) {
3    while (x <= i) ;
4    y = i + 1;
5  }
```

$$i + 1 \quad i + 1$$

```
    x = i + 1;    while (x <= i) ;
while (y <= i) ; y = i + 1;
```

$$i + 1 \quad i + 1$$

Environment

- Skylake X Gold 6130 ("big" partition of the IST cluster)
- 2 hardware threads $\times$ 16 cores $\times$ 4 sockets (= 128 processors seen by OS)
- ensure variables x and y are at least 64 bytes apart (not on the same cache line)
- bind both threads on specific processors by OpenMP environment variable OMP_BIND_PROC=true
- try all combinations of threads (i.e., with p threads, measure all the $p(p-1)$ pairs) and show a matrix

Result

- (i, j) indicates the roundtrip latency (in reference clocks) between processor i and j

src	dest	latency
0	1-15	≈ 800
0	16-63	≈ 1100
0	64	≈ 110
0	65-79	≈ 450
0	80-127	≈ 1100

- a beautiful pattern emerges which is obviously telling

Result

- e.g., which processor is "close" to processor 0?
  - 64 is closest
  - 1-15 and 65-79 are close
  - 16-63 and 80-127 are farthest
- a natural interpretation
  - x and $(x + 64)$ are two hardware threads on a core
  - 0-15 (and 65-79) are the 16 physical cores (32 hwts) on a socket
  - others are on different sockets

What they imply to parallel algorithms?

- you do not want to have many threads concurrently updating the same data
- remember SpMV COO?

```
1  // assume inside #pragma omp parallel
2      ...
3  #pragma omp for
4  for (k = 0; k < A.nnz; k++) {
5    i,j,Aij = A.elems[k];
6  #pragma omp atomic
7    y[i] += Aij * x[j];
8  }
```

- `y[i] +=` may be costing 1000 cycles when its single-thread execution would take just dozens of cycles

Summary (1): latency and bandwidth

- latency of data access heavily depends on which level of caches you actually access:

    L1 (a few cycles) ≤ main memory (> 200 cycles)

- a single core bandwidth is limited by:

$$\frac{\text{cache line size} \times \text{LFB size}}{\text{latency}}$$

- for main memory, it's much lower than what you see in the spec

- max bandwidth is attainable only with multiple cores

Summary (2): bandwidth differs by access patterns

- $$\text{bandwidth} = \frac{\text{line size} \times \text{number of accesses in flight}}{\text{latency}}$$

- bandwidth heavily depends on the number of in-flight accesses, which depend on *access patterns*
  - random address pointer chasing
  - random but independent addresses
  - sequential

Common misunderstanding

- pointer chasing is always bad
    - not when data fit in L1 (perhaps L2) cache
    - not when accessed addresses are sequential
    - not when you manage to chase many pointer chains
- random access is always worse than sequential access
    - not so much when an element ≈ cache size

Summary (3): inter processor communication

- cores communicate as a side effect of memory accesses (cache misses)
- it is natually as expensive as L2/L3 misses (or more), depending on whom you communicate with
- shared memory is nice, but you cannot forget the cost