

Neural Networks Basics

Kenjiro Taura

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

What is machine learning?

- input: a set of *training data set*

$$D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$$

- each x_i is normally a real vector (i.e. many real values)
- each t_i is a real value (regression), 0/1 (binary classification), a discrete value (multi-class classification), etc., depending on the task

What is machine learning?

- **input:** a set of *training data set*

$$D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$$

- each x_i is normally a real vector (i.e. many real values)
- each t_i is a real value (regression), 0/1 (binary classification), a discrete value (multi-class classification), etc., depending on the task
- **goal:** a supervised machine learning tries to find a function f that “matches” training data well. i.e.

$$f(x_i) \approx t_i \text{ for } (x_i, t_i) \in D$$

- put formally, find f that minimizes an *error* or a *loss*:

$$L(f; D) \equiv \sum_{(x_i, t_i) \in D} \text{err}(f(x_i), t_i),$$

where $\text{err}(y_i, t_i)$ is a function that measures an “error” or a “distance” between the predicted output and the true value

Machine learning as an optimization problem

- finding a good function from the space of *literally all* possible functions is neither easy nor meaningful
- we thus normally fix a search space of functions (\mathcal{F}) to a fixed expression parameterized by w and find a good function $f_w \in \mathcal{F}$ (*parametric models*)

Machine learning as an optimization problem

- finding a good function from the space of *literally all* possible functions is neither easy nor meaningful
- we thus normally fix a search space of functions (\mathcal{F}) to a fixed expression parameterized by w and find a good function $f_w \in \mathcal{F}$ (*parametric models*)
- the task is then to find the value of w that minimizes the loss:

$$L(w; D) \equiv \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i)$$

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

A simple example (linear regression)

- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a real value
 - t_i : a real value

A simple example (linear regression)

- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a real value
 - t_i : a real value
- let the search space be a set of polynomials of degree ≤ 2 . a function is then parameterized by $w = (w_0 \ w_1 \ w_2)$. i.e.

$$f_w(x) \equiv w_2 x^2 + w_1 x + w_0$$

A simple example (linear regression)

- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a real value
 - t_i : a real value
- let the search space be a set of polynomials of degree ≤ 2 . a function is then parameterized by $w = (w_0 \ w_1 \ w_2)$. i.e.

$$f_w(x) \equiv w_2 x^2 + w_1 x + w_0$$

- let the error function be a simple square distance:

$$\text{err}(y, t) \equiv (y - t)^2$$

A simple example (linear regression)

- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a real value
 - t_i : a real value
- let the search space be a set of polynomials of degree ≤ 2 . a function is then parameterized by $w = (w_0 \ w_1 \ w_2)$. i.e.

$$f_w(x) \equiv w_2 x^2 + w_1 x + w_0$$

- let the error function be a simple square distance:

$$\text{err}(y, t) \equiv (y - t)^2$$

- the task is to find $w = (w_0, w_1, w_2)$ that minimizes:

$$L(w; D) = \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i) = \sum_{(x_i, t_i) \in D} (w_2 x_i^2 + w_1 x_i + w_0 - t_i)^2$$

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action


A more realistic example: digit recognition

- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a vector of pixel values of an image:
 - t_i : the class of x_i ($t_i \in \{0, 1, \dots, 9\}$)

A more realistic example: digit recognition


- training data $D = \{ (x_i, t_i) \mid i = 0, 1, \dots \}$
 - x_i : a vector of pixel values of an image:
 - t_i : the class of x_i ($t_i \in \{0, 1, \dots, 9\}$)
- the search space: the following composition parameterized by three matrices W_0 and W_1

$$f_{W_0, W_1}(x) \equiv \text{softmax}(W_1 \text{maxpool}(\text{ReLU}(W_0 * x)))$$



A handwritten digits recognition


- the output $y = f_{W_0, W_1}(x)$ is a 10-vector representing probabilities that x belongs to each of the ten classes



A handwritten digits recognition

- the output $y = f_{W_0, W_1}(x)$ is a 10-vector representing probabilities that x belongs to each of the ten classes
- a loss function is *negative log-likelihood* commonly used in multiclass classifications

$$\text{err}(y, t) = \text{NLL}(y, t) \equiv -\log y_t$$




A handwritten digits recognition

- the output $y = f_{W_0, W_1}(x)$ is a 10-vector representing probabilities that x belongs to each of the ten classes
- a loss function is *negative log-likelihood* commonly used in multiclass classifications

$$\text{err}(y, t) = \text{NLL}(y, t) \equiv -\log y_t$$

- the task is to find W_0 and W_1 that minimize:

$$\begin{aligned} & L(W_0, W_1; D) \\ = & \sum_{(x_i, t_i) \in D} \text{NLL}(f_{W_0, W_1}(x_i), t_i) \\ = & \sum_{(x_i, t_i) \in D} \text{NLL}(\text{softmax}(W_1 \text{maxpool}(\text{ReLU}(W_0 * x))), t_i) \end{aligned}$$



Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

How to find the minimizing parameter?

- it boils down to minimizing a function that takes *lots of* parameters w

$$L(w; D) = \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i),$$

- we compute the derivative of L with respect to w and move w to its opposite direction (*gradient descent; GD*)

$$w = w - \eta^t \frac{\partial L}{\partial w}$$

(η : a small value controlling a learning rate)

- repeat this until $L(w; D)$ converges

Why GD works

- recall

$$L(w + \Delta w; D) \approx L(w; D) + \frac{\partial L}{\partial w} \Delta w$$

- so, by moving w slightly to the direction of gradient (i.e., $\Delta w = -\eta \frac{\partial L}{\partial w}$ for small η),

$$\begin{aligned} L(w - \eta \frac{\partial L}{\partial w}; D) &\approx L(w; D) - \eta \frac{\partial L}{\partial w} \frac{\partial L}{\partial w} \\ &< L(w; D) \end{aligned}$$

L will decrease

A linear regression example

- recall that in the linear regression example:

$$L(w; D) = \sum_{(x_i, t_i) \in D} (w_2 x_i^2 + w_1 x_i + w_0 - t_i)^2$$

A linear regression example

- recall that in the linear regression example:

$$L(w; D) = \sum_{(x_i, t_i) \in D} (w_2 x_i^2 + w_1 x_i + w_0 - t_i)^2$$

- differentiate L by $w = {}^t(w_0 \ w_1 \ w_2)$ to get:

$$\frac{\partial L}{\partial w} = \sum_{(x_i, t_i) \in D} 2(w_2 x_i^2 + w_1 x_i + w_0 - t_i) (1 \ x_i \ x_i^2)$$

(remark: we used [a chain rule](#))

A linear regression example

- recall that in the linear regression example:

$$L(w; D) = \sum_{(x_i, t_i) \in D} (w_2 x_i^2 + w_1 x_i + w_0 - t_i)^2$$

- differentiate L by $w = {}^t(w_0 \ w_1 \ w_2)$ to get:

$$\frac{\partial L}{\partial w} = \sum_{(x_i, t_i) \in D} 2(w_2 x_i^2 + w_1 x_i + w_0 - t_i) \begin{pmatrix} 1 \\ x_i \\ x_i^2 \end{pmatrix}$$

(remark: we used [a chain rule](#))

- so you repeat:

$$w = w - \eta \sum_{(x_i, t_i) \in D} 2(w_2 x_i^2 + w_1 x_i + w_0 - t_i) \begin{pmatrix} 1 \\ x_i \\ x_i^2 \end{pmatrix}$$

until $L(w; D)$ converges

A problem of the gradient descent

- the loss function we want to minimize is normally a summation over *all* training data:

$$L(w; D) = \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i)$$

A problem of the gradient descent

- the loss function we want to minimize is normally a summation over *all* training data:

$$L(w; D) = \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i)$$

- the gradient descent method just described:
 - computes $\frac{\partial}{\partial w} \text{err}(f_w(x_i), t_i)$ for each training data $(x_i, t_i) \in D$, *with the current value of w*
 - sum them over *whole data set* and then update w

A problem of the gradient descent

- the loss function we want to minimize is normally a summation over *all* training data:

$$L(w; D) = \sum_{(x_i, t_i) \in D} \text{err}(f_w(x_i), t_i)$$

- the gradient descent method just described:
 - computes $\frac{\partial}{\partial w} \text{err}(f_w(x_i), t_i)$ for each training data $(x_i, t_i) \in D$, *with the current value of w*
 - sum them over *whole data set* and then update w
- it is commonly observed that the convergence becomes faster when we update w more “incrementally” \rightarrow *Stochastic Gradient Descent (SGD)*

Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

repeat:

- 1 randomly draw a *subset* of training data D' (a mini batch; $D' \subset D$)

repeat:

- ① randomly draw a *subset* of training data D' (a mini batch; $D' \subset D$)
- ② compute the gradient of loss *over the mini batch*

$$\frac{\partial L(w; D')}{\partial w} = \sum_{(x_i, t_i) \in D'} \frac{\partial}{\partial w} \text{err}(f_w(x_i), t_i)$$

repeat:

- 1 randomly draw a *subset* of training data D' (a mini batch; $D' \subset D$)
- 2 compute the gradient of loss *over the mini batch*

$$\frac{\partial L(w; D')}{\partial w} = \sum_{(x_i, t_i) \in D'} \frac{\partial}{\partial w} \text{err}(f_w(x_i), t_i)$$

- 3 update w

$$w = w - \eta^t \frac{\partial L(w; D')}{\partial w}$$

- 4 “update sooner rather than later”

Computing the gradients

- in neural networks, a function is a composition of many stages each represented by a lot of parameters

$$x_1 = f_1(w_1; x)$$


$$x_2 = f_2(w_2; x_1)$$

...

$$y = f_n(w_n; x_n)$$

$$\ell = \text{err}(y, t)$$


- we need to differentiate ℓ by w_1, \dots, w_n



The digit recognition example

$$\begin{aligned}x_1 &= W_0 * x \\x_2 &= \text{ReLU}(x_1) \\x_3 &= \text{maxpool}(x_2) \\x_4 &= W_1 x_3 \\y &= \text{softmax}(x_4) \\\ell &= \text{NLL}(y, t)\end{aligned}$$

you need to differentiate ℓ by W_0 and W_1



Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

Differentiating multivariable functions

- $x = {}^t(x_0 \cdots x_{n-1}) \in R^n$ (a column vector)
- $f(x)$: a scalar
- **definition:** the gradient of f with respect to x , written $\frac{\partial f}{\partial x}$, is a row n -vector s.t.

$$\begin{aligned}\Delta f &\equiv f(x + \Delta x) - f(x) \\ &\approx \frac{\partial f}{\partial x} \Delta x \\ &= \sum_{i=0}^{n-1} \left(\frac{\partial f}{\partial x} \right)_i \Delta x_i\end{aligned}$$

- when it exists,

$$\frac{\partial f}{\partial x} = \left(\frac{\partial f}{\partial x_0} \cdots \frac{\partial f}{\partial x_{n-1}} \right),$$


so

$$\Delta f \approx \sum_{i=0}^{n-1} \frac{\partial f}{\partial x_i} \Delta x_i$$


The Chain Rule

- consider a function f that depends on $y = (y_0, \dots, y_{m-1}) \in R^m$, each of which in turn depends on $x = (x_0, \dots, x_{n-1}) \in R^n$
- the chain rule (math textbook version):

$$\frac{\partial f}{\partial x_i} = \sum_{0 \leq j < m} \frac{\partial f}{\partial y_j} \frac{\partial y_j}{\partial x_i} \quad (0 \leq i < n)$$



The Chain Rule : intuition



- say you increase an input variable x_i by Δx_i , each y_j will increase by

$$\approx \frac{\partial y_j}{\partial x_i} \Delta x_i,$$


which will contribute to increasing the final output (f) by

$$\approx \frac{\partial f}{\partial y_j} \frac{\partial y_j}{\partial x_i} \Delta x_i$$

Chain Rule

- master the following “index-free” version for neural network
- x, y : a scalar (a single component in a vector/matrix/high dimensional array)
- the chain rule (ML practitioner's version):

$$\frac{\partial f}{\partial x} = \sum_{\text{all variables } y \text{ that } x \text{ directly affects}} \frac{\partial f}{\partial y} \frac{\partial y}{\partial x}$$



Chain Rule and “Back Propagation”

- Chain rule allows you to compute


$$\frac{\partial L}{\partial x},$$

the derivative of the loss with respect to a variable, from

$$\frac{\partial L}{\partial y},$$

the derivatives of the loss with respect to upstream variables

$$\frac{\partial L}{\partial x} = \sum_{\text{all variables } y \text{ a step ahead of } x} \frac{\partial L}{\partial y} \frac{\partial y}{\partial x}$$



Contents

- 1 What is machine learning?
 - A simple linear regression
 - A handwritten digit recognition
- 2 Training
 - A simple gradient descent
 - Stochastic gradient descent
- 3 Chain Rule
- 4 Back Propagation in Action

Component functions


we use the following functions

- $\text{Convolution}(W; x)$: applies a linear filter
- $\text{Linear}(W; x)$: multiplies x by W
- $\text{ReLU}(x)$: zero negative values
- $\text{maxpool}(x)$: replaces each 2x2 patch with 1x1
- $\text{dropout}(x)$: probabilistically zeros some values
- $\text{softmax}(x)$: normalizes x and amplifies large values
- $\text{NLL}(x, t)$: negative log-likelihood

we summarize their definitions and their derivatives


Convolution

- it takes
 - an image = 2D pixels \times a number of channels
 - a “filter” or a “kernel”, which is essentially a small image and slides the filter over all pixels of the input and takes the local inner product at each pixel
- an illustration of a single channel 2D convolution (imagine a grayscale image)



Convolution (a single channel version)

- $W_{i,j}$: a filter ($0 \leq i < K$, $0 \leq j < K$)
- b : bias
- $x_{i,j}$: an input image ($0 \leq i < H$, $0 \leq j < W$)
- $y_{i,j}$: an output image ($0 \leq i < H - K + 1$, $0 \leq j < W - K + 1$)



$$\forall i, j \quad y_{i,j} = \sum_{0 \leq i' < K, 0 \leq j' < K} w_{i',j'} x_{i+i',j+j'} + b$$

Convolution (multiple channels version)

- say input has IC channels and output OC channels
- $W_{oc,ic,i,j}$: filter ($0 \leq ic < IC$, $0 \leq oc < OC$)
- b_{oc} : bias ($0 \leq oc < OC$)
- $x_{ic,i,j}$: an input image
- $y_{oc,i,j}$: an output image

$$\forall oc, i, j \quad y_{oc,i,j} = \sum_{ic,i',j'} w_{oc,ic,i',j'} x_{ic,i+i',j+j'} + b_{oc}$$

- the actual code does this for each sample in a batch

$$\forall s, oc, i, j \quad y_{s,oc,i,j} = \sum_{ic,i',j'} w_{oc,ic,i',j'} x_{s,ic,i+i',j+j'} + b_{oc}$$

Convolution (Back propagation 1)

- $\frac{\partial L}{\partial x}$

$$\begin{aligned}\frac{\partial L}{\partial x_{s,ic,i+i',j+j'}} &= \sum_{s',oc,i,j} \frac{\partial L}{\partial y_{s',oc,i,j}} \frac{\partial y_{s',oc,i,j}}{\partial x_{s,ic,i+i',j+j'}} \\ &= \sum_{oc,i,j} \frac{\partial L}{\partial y_{s,oc,i,j}} w_{oc,ic,i',j'}\end{aligned}$$

Convolution (Back propagation 2)

- $\frac{\partial L}{\partial w}$

$$\begin{aligned}\frac{\partial L}{\partial w_{oc,ic,i',j'}} &= \sum_{s,oc',i,j} \frac{\partial L}{\partial y_{s,oc',i,j}} \frac{\partial y_{s,oc',i,j}}{\partial w_{oc,ic,i',j'}} \\ &= \sum_{s,i,j} \frac{\partial L}{\partial y_{s,oc,i,j}} x_{s,ic,i+i',j+j'}\end{aligned}$$

- $\frac{\partial L}{\partial b}$

$$\begin{aligned}\frac{\partial L}{\partial b_{oc}} &= \sum_{s,oc',i,j} \frac{\partial L}{\partial b_{oc}} \frac{\partial y_{s,oc',i,j}}{\partial b_{oc}} \\ &= \sum_{s,i,j} \frac{\partial L}{\partial y_{s,oc,i,j}}\end{aligned}$$

Linear (a.k.a. Fully Connected Layer)

- **definition:**

$$\begin{aligned} y = \text{Linear}(W; x) &\equiv Wx + b \\ \forall i \quad y_i &= \sum_j W_{ij}x_j + b_i \end{aligned}$$

Linear (Back Propagation 1)

- $\frac{\partial L}{\partial x}$

$$\begin{aligned}\frac{\partial L}{\partial x_j} &= \sum_{i'} \frac{\partial L}{\partial y_{i'}} \frac{\partial y_{i'}}{\partial x_j} \\ &= \sum_{i'} \frac{\partial L}{\partial y_{i'}} w_{i'j}\end{aligned}$$

Linear (Back Propagation 2)

- $\frac{\partial L}{\partial W}$

$$\begin{aligned}\frac{\partial L}{\partial W_{ij}} &= \sum_{i'} \frac{\partial L}{\partial y_{i'}} \frac{\partial y_{i'}}{\partial W_{ij}} \\ &= \frac{\partial L}{\partial y_i} x_j\end{aligned}$$

- $\frac{\partial L}{\partial b}$

$$\begin{aligned}\frac{\partial L}{\partial b_i} &= \sum_{i'} \frac{\partial L}{\partial y_{i'}} \frac{\partial y_{i'}}{\partial b_i} \\ &= \frac{\partial L}{\partial y_i}\end{aligned}$$


ReLU

- **definition (scalar ReLU):** for $x \in R$, define

$$\text{relu}(x) \equiv \max(x, 0)$$

- **derivatives of relu:** for $y = \text{relu}(x)$,

$$\frac{\partial y}{\partial x} = \begin{cases} 1 & (x > 0) \\ 0 & (x \leq 0) \end{cases} = \max(\text{sign}(x), 0)$$



- **definition (vector ReLU):** for a vector $x \in R^n$, define ReLU as the application of relu to each component

$$\text{ReLU}(x) \equiv \begin{pmatrix} \text{relu}(x_0) \\ \vdots \\ \text{relu}(x_{n-1}) \end{pmatrix}$$

- **derivatives of ReLU:**

$$\frac{\partial y_j}{\partial x_i} = \begin{cases} \max(\text{sign}(x_i), 0) & (i = j) \\ 0 & (i \neq j) \end{cases}$$

- back propagation:

$$\begin{aligned}\frac{\partial L}{\partial x_j} &= \sum_i \frac{\partial L}{\partial y_i} \frac{\partial y_i}{\partial x_j} \\ &= \frac{\partial L}{\partial y_j} \frac{\partial y_j}{\partial x_j} \\ &= \begin{cases} \frac{\partial L}{\partial y_j} & (x_j \geq 0) \\ 0 & (x_j < 0) \end{cases}\end{aligned}$$

softmax

- **definition:** for $x \in R^n$

$$y = \text{softmax}(x) \equiv \frac{1}{\sum_{i=0}^{n-1} \exp(x_i)} \begin{pmatrix} \exp(x_0) \\ \vdots \\ \exp(x_{n-1}) \end{pmatrix}$$

it is a vector whose:

- each component > 0 ,
- sum of all components = 1
- largest component “dominates”

$$\begin{aligned} y &= \log(\text{softmax}(x)) \\ &= \begin{pmatrix} x_0 - \log \sum_{i=0}^{n-1} \exp(x_i) \\ \vdots \\ x_{n-1} - \log \sum_{i=0}^{n-1} \exp(x_i) \end{pmatrix} \end{aligned}$$

- (recall

$$\text{softmax}(x) \equiv \frac{1}{\sum_{i=0}^{n-1} \exp(x_i)} \begin{pmatrix} \exp(x_0) \\ \vdots \\ \exp(x_{n-1}) \end{pmatrix}$$

)

- **definition:**

- x : n -vector
- t : true class of the data

$$\text{NLL}(x, t) \equiv -\log x_t$$

- thus,

$$\begin{aligned} y &= \text{NLL}(\text{softmax}(x), t) \\ &= -x_t + \log \sum_{i=0}^{n-1} \exp(x_i) \end{aligned}$$

NLL softmax (Back propagation)



$$\begin{aligned}\frac{\partial L}{\partial x_i} &= \frac{\partial L}{\partial y} \frac{\partial y}{\partial x_i} \\ &= \begin{cases} \frac{\partial L}{\partial y} \left(-1 + \frac{\exp(x_i)}{\sum_{i=0}^{n-1} \exp(x_i)} \right) & (i = t) \\ \frac{\partial L}{\partial y} \frac{\exp(x_i)}{\sum_{i=0}^{n-1} \exp(x_i)} & (i \neq t) \end{cases} \\ &= \begin{cases} \frac{\partial L}{\partial y} (-1 + \text{softmax}(x_i)) & (i = t) \\ \frac{\partial L}{\partial y} \text{softmax}(x_i) & (i \neq t) \end{cases}\end{aligned}$$

Note: why NLL softmax?

- recall that for n -way classification, the output of $p = \text{softmax}(\dots)$ is an n -vector
- p_i is meant to be the *probability* that a particular sample belongs to the class i
- for that purpose, a loss function could be any function that decreases with p_t (something as simple as $-p_t$), where t is the true label of the particular sample
- we instead use $\text{NLL}(p, t) = -\log p_t$. why?

Note: why NLL log softmax?

- this is because,

- ① the goal is to maximize the joint probability of the entire data, which is the *product* of probabilities of individual samples:

$$\prod_k p_{t_k},$$

where t_k is the true label of sample k , and

- ② the loss over a mini-batch is the *sum* of losses of individual samples
- they can be reconciled by setting the loss function to $-\log p_t$

$$\sum_k (-\log p_{t_k}) = -\log (\prod_k p_{t_k})$$